

YouTube Data Analyzer

Team Kongqueror

Ross Kugler, Benjamin Bordon, Huy (Harry) Ky

Problem Statement

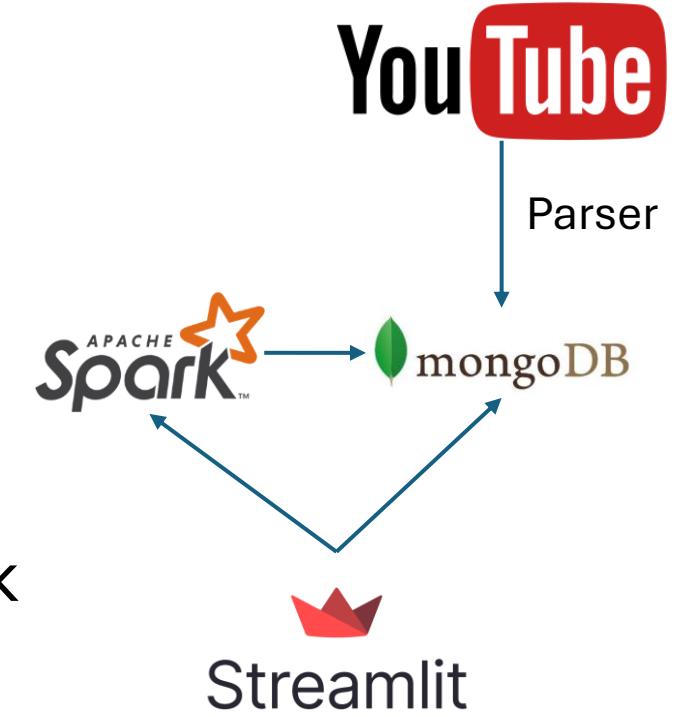
- Implement a Youtube data analyzer supported by NoSQL database, Hadoop MapReduce, and Spark GraphX/GraphFrame. The analyzer provides basic data analytics functions to Youtube media datasets.
- Efficiently storing, processing, and analyzing large-scale YouTube video network data to extract meaningful insights

Team

- 3 senior Software Engineering students at WSU Everett
 - Ross
 - Ben
 - Harry

Architecture

- Parse the crawls from the dataset
- Insert into MongoDB
- Spark connects to MongoDB via connector
- Python GUI application using Streamlit framework
 - Connects to both MongoDB and Spark
 - Loads precomputed algorithm results from MongoDB
 - Can dynamically run Spark algorithms based on user queries
- Spark Algorithms written in Python with PySpark
- Apache Spark
 - Run in Standalone mode require at least 1 worker with 4GB of memory



Dataset Parsing

- Initially took 2 crawls from the initial dataset:
<https://netsg.cs.sfu.ca/youtubedata/>
- Parse through individual crawls
- Clean to remove malformed/incomplete data
- Create indexes
- Insert into 4 MongoDB collections
 - crawls
 - edges
 - video_snapshots
 - videos
- Raw Data (245 Mb) -> Inserted Data (1.77 GiB)

Spark Cluster

- Spark 3.5.7
- Runs in standalone mode in Windows
- Number of workers based on local resources (# of CPU cores and memory)
- We reserve 4 cores for OS
- Optimized the remaining CPU cores for the Spark worker nodes.
 - Prioritize higher memory per worker than multiple workers with lower memory.

Algorithms & Queries

- Network Aggregation

A high-level summary of the entire video interaction graph, capturing its scale and connectivity.

- Categorized Video Stats

Performance metrics broken down by video categories (entertainment, education, comedy, etc.)

- Top-K

- # of videos in a Category
- By view
- By rating
- By range

- Subgraph Patterns

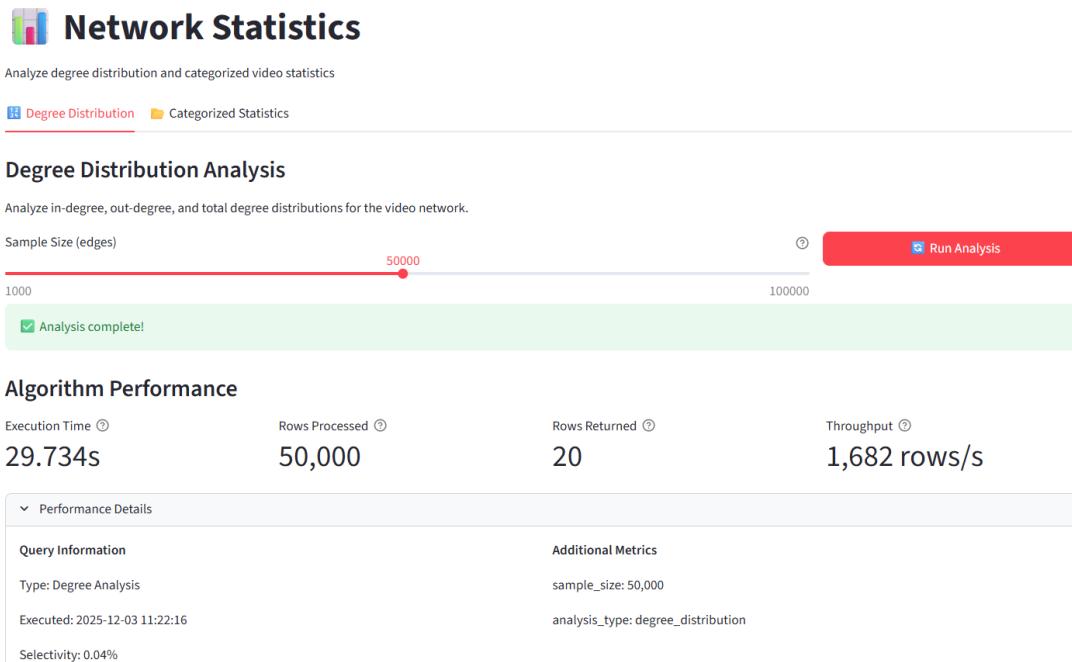
Detects common structural motifs—hubs, chains, cliques, etc.—to uncover hidden audience or influencer dynamics.

- Pagerank

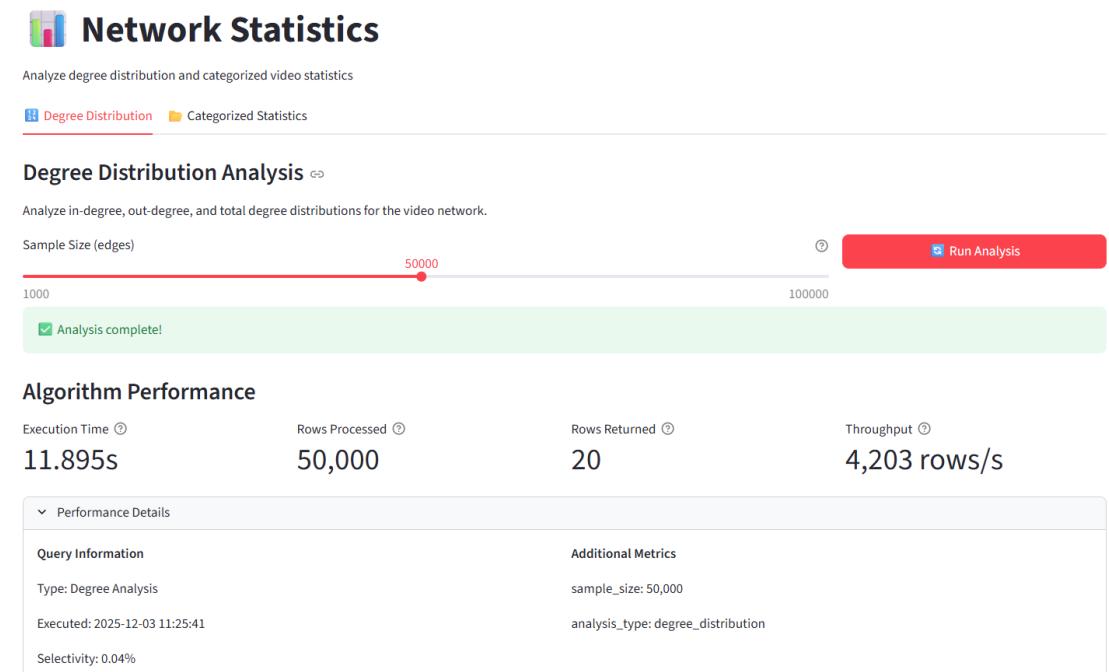
Ranks videos by influence, based on the quality and quantity of connections from other important videos.

Performance Results

- Live performance results provided in app



Cold start (first run)



After caching (second run)

Data Insights

- *Keep in mind this is limited to the small sample of the dataset we used:
 - Music, entertainment are the top category of all time.
 - GRANHERMANO8 has most videos – Spanish Reality TV show, followed by CBS
 - Education was surprisingly one of the top categories, even though online schooling was likely not that popular in 2007-2008
 - Highest viewed video is over 6 million views, compared to current day (December 2025) data of 16 Billion views.

https://en.wikipedia.org/wiki/List_of_most-viewed_YouTube_videos

Other Insights

- The dataset we worked with was old (2007-2008)
 - Many videos are unavailable – either violated terms or taken down by uploader
 - Would've been more beneficial to do our own crawling with current data
- “una” video category from the crawl dataset?

Lessons Learned

- Probably not do as much indexing as we did (unless we are going full-scale), because it made the dataset expand
 - If we had time and resources, actually use Docker containers and deploy the app, with Spark nodes on different machines

```
PS C:\Users\rossk> mongosh mongodb://localhost:27017/youtube_analytics
Current Mongosh Log ID: 69309d0ae55456941e1e2620
Connecting to:      mongodb://localhost:27017/youtube_analytics?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+2.5.10
Using MongoDB:     8.2.0
Using Mongosh:    2.5.10

For mongosh info see: https://www.mongodb.com/docs/mongodb-shell/

-----
The server generated these startup warnings when booting
2025-11-29T20:47:11.741-08:00: Access control is not enabled for the database. Read and write access to data and configuration is unrestricted
-----
rs0 [direct: primary] youtube_analytics> db.stats()
{
  db: 'youtube_analytics',
  collections: Long('4'),
  views: Long('0'),
  objects: Long('19102596'),
  avgObjSize: 92.54287867470997,
  dataSize: 1767889224,
  storageSize: 4886446688,
  indexes: Long('21'),
  indexSize: 1332555776,
  totalSize: 1821200384,
  scaleFactor: Long('1'),
  fsUsedSize: 520443277312,
  fsTotalSize: 998811254784,
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1764793607, i: 1 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAA='),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1764793607, i: 1 })
}
```



Thank you

- Questions?